

Assignment7-Week10-SentimentAnalysis

Ross Boehme

2023-04-02

Assignment 7

- In Text Mining with R, Chapter 2 looks at Sentiment Analysis. In this assignment, you should start by getting the primary example code from chapter 2 working in an R Markdown document. You're then asked to extend the code in two ways:
1. Work with a different corpus of your choosing
 2. Incorporate at least one additional sentiment lexicon (possibly from another R package that you've found through research).

Initial Imports and Citations

```
library(tidytext)
```

```
## Warning: package 'tidytext' was built under R version 4.2.3
```

```
library(janeaustenr)
```

```
## Warning: package 'janeaustenr' was built under R version 4.2.3
```

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.4.0    v purrr   1.0.1
## v tibble  3.1.8    v dplyr  1.1.0
## v tidyr   1.3.0    v stringr 1.5.0
## v readr   2.1.4    v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(textdata)
```

```
## Warning: package 'textdata' was built under R version 4.2.3
```

Please note the base code for this project was taken from “Text Mining with R: A Tidy Approach” by Julia Silge and David Robinson. Sebastopol, CA: O’Reilly Media, 2017. ISBN 978-1-491-98165-8. Code below in named chunks “silge-robinson” and “silge-robinson2”.

This first R chunk code shows how to “mine” words from Jane Austen’s bibliography, with each observation in the dataframe a unique Book, Chapter, Line Number, and Word combination.

```
tidy_books <- austen_books() %>%
  group_by(book) %>%
  mutate(
    linenumber = row_number(),
    chapter = cumsum(str_detect(text,
                                regex("^chapter [\\divxlc]",
                                       ignore_case = TRUE)))) %>%
  ungroup() %>%
  unnest_tokens(word, text)
```

This second R chunk shows how to load three sentiment lexicons: afinn, Bing, and NRC which I’d also like to cite (linked below). - afinn - Bing - nrc

```
afinn <- get_sentiments('afinn')
bing <- get_sentiments('bing')
nrc <- get_sentiments('nrc')
```

I will use this base as inspiration for my own analysis.

Extending Base Code With Same Corpus

I will look at the full available Jane Austen corpus to extend this code, evaluating which of the author’s books have the highest percentage of neutral, positive, and negative words.

Showing word count by book. Sense & Sensibility is the longest book. Persuasion is the shortest.

```
austen_word_counts <- as.data.frame(table(tidy_books$book))
colnames(austen_word_counts) <- c("book", "total_word_count")
austen_word_counts
```

##	book	total_word_count
## 1	Sense & Sensibility	119957
## 2	Pride & Prejudice	122204
## 3	Mansfield Park	160460
## 4	Emma	160996
## 5	Northanger Abbey	77780
## 6	Persuasion	83658

Taking positive and negative word lexicon from nrc which I’ll use to evaluate Jane Austen’s corpus. Merging with tidy_books df created previously, then showing positive and negative word count by book.

```
nrc_pos_neg <- nrc %>%
  filter(sentiment == 'negative' | sentiment == 'positive')

austen_sentiment <- tidy_books %>%
```

```
inner_join(nrc_pos_neg) %>%
group_by(book,sentiment) %>%
summarise(total_count=n(),.groups = 'drop') %>%
as.data.frame() %>%
pivot_wider(names_from = sentiment,
            values_from = total_count)
```

```
## Joining with 'by = join_by(word)'
```

```
## Warning in inner_join(., nrc_pos_neg): Each row in 'x' is expected to match at most 1 row in 'y'.
## i Row 251 of 'x' matches multiple rows.
## i If multiple matches are expected, set 'multiple = "all"' to silence this
## warning.
```

```
colnames(austen_sentiment) <- c("book","negative_words","positive_words")
```

```
austen_sentiment
```

```
## # A tibble: 6 x 3
##   book                negative_words positive_words
##   <fct>                <int>          <int>
## 1 Sense & Sensibility      3994            7314
## 2 Pride & Prejudice        3630            7273
## 3 Mansfield Park          4694            9269
## 4 Emma                    4438            9262
## 5 Northanger Abbey        2456            4520
## 6 Persuasion              2226            5248
```

Merging total word count df with positive and negative word count df to assess % of all words which were positive or negative by book.

```
austen_combined <- austen_word_counts %>%
  inner_join(austen_sentiment) %>%
  transform(perc_neg = scales::percent(negative_words / total_word_count),
            perc_pos = scales::percent(positive_words / total_word_count),
            neutral_words = total_word_count - negative_words - positive_words)
```

```
## Joining with 'by = join_by(book)'
```

```
austen_combined <- austen_combined %>%
  transform(perc_neutr = scales::percent(neutral_words / total_word_count))
```

```
austen_final <- austen_combined[,c('book','total_word_count','positive_words','negative_words','neutral_words')]
```

```
austen_final
```

```
##           book total_word_count positive_words negative_words
## 1 Sense & Sensibility      119957            7314          3994
## 2 Pride & Prejudice        122204            7273          3630
## 3 Mansfield Park          160460            9269          4694
```

## 4	Emma	160996	9262	4438
## 5	Northanger Abbey	77780	4520	2456
## 6	Persuasion	83658	5248	2226
##	neutral_words	perc_pos	perc_neg	perc_neutr
## 1	108649	6.097%	3.330%	90.573%
## 2	111301	5.952%	2.970%	91.078%
## 3	146497	5.777%	2.925%	91.298%
## 4	147296	5.753%	2.757%	91.490%
## 5	70804	5.811%	3.158%	91.031%
## 6	76184	6.273%	2.661%	91.066%

The book with the highest percentage of positive words was: Persuasion. It also had the lowest percentage of negative words.

```
austen_final %>% arrange(desc(perc_pos))
```

##	book	total_word_count	positive_words	negative_words
## 1	Persuasion	83658	5248	2226
## 2	Sense & Sensibility	119957	7314	3994
## 3	Pride & Prejudice	122204	7273	3630
## 4	Northanger Abbey	77780	4520	2456
## 5	Mansfield Park	160460	9269	4694
## 6	Emma	160996	9262	4438
##	neutral_words	perc_pos	perc_neg	perc_neutr
## 1	76184	6.273%	2.661%	91.066%
## 2	108649	6.097%	3.330%	90.573%
## 3	111301	5.952%	2.970%	91.078%
## 4	70804	5.811%	3.158%	91.031%
## 5	146497	5.777%	2.925%	91.298%
## 6	147296	5.753%	2.757%	91.490%

The book with the highest percentage of negative words was “Sense & Sensibility” which is interesting because it also had the second highest percentage of positive words. It therefore had fewer “neutral” words, making it a more emotionally charged book than the others.

```
austen_final %>% arrange(desc(perc_neg))
```

##	book	total_word_count	positive_words	negative_words
## 1	Sense & Sensibility	119957	7314	3994
## 2	Northanger Abbey	77780	4520	2456
## 3	Pride & Prejudice	122204	7273	3630
## 4	Mansfield Park	160460	9269	4694
## 5	Emma	160996	9262	4438
## 6	Persuasion	83658	5248	2226
##	neutral_words	perc_pos	perc_neg	perc_neutr
## 1	108649	6.097%	3.330%	90.573%
## 2	70804	5.811%	3.158%	91.031%
## 3	111301	5.952%	2.970%	91.078%
## 4	146497	5.777%	2.925%	91.298%
## 5	147296	5.753%	2.757%	91.490%
## 6	76184	6.273%	2.661%	91.066%

Additional Sentiment Analysis

Per the assignment description, I will now perform an additional sentiment analysis using a different corpus and different lexicon than previously mentioned.

For the lexicon I will use SentimentAnalysis.R package description linked [here](#).

For the corpus I will use the “friends” package from R where each observation is a piece of speech said by a character in the TV show Friends.

```
library(SentimentAnalysis)
```

```
## Warning: package 'SentimentAnalysis' was built under R version 4.2.3
```

```
##
```

```
## Attaching package: 'SentimentAnalysis'
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
##      write
```

```
library(friends)
```

```
## Warning: package 'friends' was built under R version 4.2.3
```

```
friends_lines <- friends
```

Data Cleaning and Exploratory Data Analysis As the value count generated below shows, there are 699 characters who have speaking lines over the course of Friends. To ensure an adequate sample size, I'll only look at the six main characters who account for the vast majority of lines: Monica Geller, Rachel Green, Ross Geller, Chandler Bing, Joey Tribbiani, and Phoebe Buffay.

Per this dataframe, the characters with the most lines are: Rachel Green (9312), Ross Geller (9157), and Chandler Bing (8465).

```
length(table(friends_lines$speaker))
```

```
## [1] 699
```

```
talkers <- as.data.frame(table(friends_lines$speaker))
top_10_talkers <- talkers %>%
  slice_max(order_by = Freq, n = 10)
names(top_10_talkers) <- c('character', 'speaking_lines')
top_10_talkers
```

```
##      character speaking_lines
## 1    Rachel Green          9312
## 2     Ross Geller          9157
## 3   Chandler Bing          8465
## 4   Monica Geller          8441
## 5   Joey Tribbiani          8171
```

```
## 6      Phoebe Buffay          7501
## 7  Scene Directions          6063
## 8           #ALL#           347
## 9      Mike Hannigan          330
## 10     Richard Burke          281
```

Only top 6 characters in terms of lines spoken.

```
friends_lines <- friends_lines %>%
  filter(speaker %in% c("Monica Geller", "Rachel Green", "Ross Geller", "Chandler Bing", "Joey Tribbiani"))
```

Splitting each row into multiple rows, where each word is its own row (to prepare for sentiment analysis). Delimiter between words is space " ".

```
friends_words <- friends_lines %>%
  separate_rows(text, sep = " ")
```

Showing which friends characters spoke the most, in terms of lines, words, and words per line. Monica has the fewest words per line (9.8) while Phoebe has the most (10.9). Rachel and Ross have the most lines overall (9312 and 9157) which drives their top total word count (97,633 and 95,561) among all characters.

```
friends_word_count <- friends_words %>%
  group_by(speaker) %>%
  summarise(total_word_count=n(),
            .groups = 'drop')

top_talkers_lines <- top_10_talkers %>%
  filter(character %in% c("Monica Geller", "Rachel Green", "Ross Geller", "Chandler Bing", "Joey Tribbiani"))

friends_summary <- left_join(friends_word_count, top_talkers_lines, by=c('speaker'='character'))

friends_summary <- friends_summary %>%
  transform(words_per_line = round((total_word_count / speaking_lines),2))

friends_summary
```

```
##           speaker total_word_count speaking_lines words_per_line
## 1 Chandler Bing      86547           8465         10.22
## 2 Joey Tribbiani      86426           8171         10.58
## 3 Monica Geller      82988           8441          9.83
## 4 Phoebe Buffay      81506           7501         10.87
## 5 Rachel Green      97633           9312         10.48
## 6 Ross Geller       95561           9157         10.44
```

Sentiment Analysis I first attempted to perform sentiment analysis using the `analyzeSentiment` function in the `SentimentAnalysis` package. However that function is too slow to work well on individual words. EG `friends_words.sentiment <- analyzeSentiment(friends_words.text)`. It's designed for smaller samples (e.g. a few paragraphs).

Therefore I'll change the `SentimentAnalysis`' word dictionary ("DictionaryGI") to a dataframe, and perform an analysis by joining the word dataframe with the sentiment dataframe, as we did with the Jane Austen data.

Preparing sentiment df.

```
data(DictionaryGI)
str(DictionaryGI)
```

```
## List of 2
## $ negative: chr [1:2005] "abandon" "abandonment" "abate" "abdicate" ...
## $ positive: chr [1:1637] "abide" "ability" "able" "abound" ...
```

```
length(DictionaryGI$positive) <- length(DictionaryGI$negative)
```

```
sa_df <- as.data.frame(DictionaryGI)
```

```
neg_words <- sa_df$negative
neg_words <- as.data.frame(neg_words)
neg_words <- neg_words %>%
  mutate(sentiment = "negative")
names(neg_words) <- c('text', 'sentiment')
```

```
pos_words <- sa_df$positive
pos_words <- as.data.frame(pos_words)
pos_words <- pos_words %>%
  mutate(sentiment="positive")
names(pos_words) <- c('text', 'sentiment')
```

```
sa_dict <- bind_rows(pos_words, neg_words)
```

```
friends_sentiment <- friends_words %>%
  inner_join(sa_dict) %>%
  group_by(speaker, sentiment) %>%
  summarise(total_count=n(), .groups = 'drop') %>%
  as.data.frame() %>%
  pivot_wider(names_from = sentiment,
              values_from = total_count)
```

```
## Joining with 'by = join_by(text)'
```

```
## Warning in inner_join(., sa_dict): Each row in 'x' is expected to match at most 1 row in 'y'.
## i Row 69 of 'x' matches multiple rows.
## i If multiple matches are expected, set 'multiple = "all"' to silence this
## warning.
```

```
colnames(friends_sentiment) <- c("speaker", "negative_words", "positive_words")
```

Showing count of positive and negative words used by Friends characters.

```
friends_sentiment
```

```
## # A tibble: 6 x 3
##   speaker      negative_words positive_words
##   <chr>          <int>          <int>
## 1 Chandler Bing      2299          3512
## 2 Joey Tribbiani     2288          3564
```

```
## 3 Monica Geller          2303          3335
## 4 Phoebe Buffay          2030          3564
## 5 Rachel Green           2477          4321
## 6 Ross Geller            2329          3822
```

Merging total Friends word count df with positive and negative word count df to assess % of all words which were positive or negative by character.

```
friends_combined <- friends_summary %>%
  inner_join(friends_sentiment) %>%
  transform(perc_neg = scales::percent(negative_words / total_word_count),
            perc_pos = scales::percent(positive_words / total_word_count),
            neutral_words = total_word_count - negative_words - positive_words)
```

```
## Joining with 'by = join_by(speaker)'
```

```
friends_combined <- friends_combined %>%
  transform(perc_neutr = scales::percent(neutral_words / total_word_count))
```

```
friends_final <- friends_combined[,c('speaker', 'total_word_count', 'speaking_lines', 'words_per_line', 'positive_words', 'negative_words', 'neutral_words', 'perc_pos', 'perc_neg', 'perc_neutr')]
```

In general, this dataframe shows the characters with a relatively narrow band of sentiment: from 4.00% positive (Ross) to 4.43% positive (Phoebe) and from 2.43% negative (Ross) to 2.78% negative (Monica). Neutrality ranged from 93.04% (Rachel) to 93.56% (Ross).

These data show Rachel as the most sentimental character, Ross as the least sentimental, Phoebe as the most positive, and Monica as the most negative. All of this aligns with my domain knowledge, therefore the sentiment analysis appears successful.

```
friends_final
```

```
##      speaker total_word_count speaking_lines words_per_line positive_words
## 1 Chandler Bing          86547          8465          10.22          3512
## 2 Joey Tribbiani          86426          8171          10.58          3564
## 3 Monica Geller          82988          8441           9.83          3335
## 4 Phoebe Buffay          81506          7501          10.87          3564
## 5 Rachel Green          97633          9312          10.48          4321
## 6 Ross Geller           95561          9157          10.44          3822
##      negative_words neutral_words perc_pos perc_neg perc_neutr
## 1           2299          80736  4.058%  2.6564%  93.286%
## 2           2288          80574  4.124%  2.6474%  93.229%
## 3           2303          77350  4.019%  2.7751%  93.206%
## 4           2030          75912  4.373%  2.4906%  93.137%
## 5           2477          90835  4.426%  2.5371%  93.037%
## 6           2329          89410  4.000%  2.4372%  93.563%
```